

SUPPLEMENTARY MATERIAL FOR “SURVEY-WEIGHTED ORDINAL CAUSAL DISCOVERY: APPLICATION TO VETERAN SOCIAL NEEDS, MENTAL HEALTH, AND WELL-BEING”

BY SOUMIK PURKAYASTHA^{1,2,a}, DAVID A. FRANK^{2,b}, LAUREN E. RUSSELL^{3,c} AND LESLIE R. M. HAUSMANN^{2,4,d}

¹Department of Biostatistics and Health Data Science, University of Pittsburgh School of Public Health, ^asoumik@pitt.edu

²Center for Healthcare Evaluation, Research, and Promotion, Veterans Affairs Pittsburgh Healthcare System, ^bDavid.Frank2@va.gov

³Office of Health Equity, Veterans Health Administration, ^cLauren.Russell3@va.gov

⁴Department of Medicine, University of Pittsburgh School of Medicine, ^dLeslie.Hausmann@va.gov

This supplement provides: §1 the `swa-oBn` algorithm in pseudocode; §2 full proofs of all theoretical results stated in the main text; §3 details of the polychoric clustering procedure; §4 simulation Study 1 supplementary metrics (edge precision and correct orientation rate); §5 simulation Study 3 supplementary tables (paired SHD gaps across $\mathcal{L}$ -regimes and sample sizes, weighted and unweighted arms); §6 simulation Study 4 (survey-weighted PC, FCI, and LiNGAM); §7 proportional odds diagnostics; §8 survey weight distribution details; §9 the composite node table linking the analytic nodes to their constituent SHEP items; §10 the complete 10×10 bootstrap edge probability matrix from the application; and §11 the weighted-versus-unweighted comparison. References to assumptions, propositions, and theorems refer to the main text unless otherwise noted.

1. The `swa-oBn` algorithm. The `swa-oBn` procedure is parameterized by the scoring choice `score_type` $\in \{\text{SINGLE-BIC}, \text{SPLIT-BIC}\}$ defined in Section 4 of the main text. Algorithm 1 formalizes the move-evaluation function used inside the greedy search; Algorithm 2 formalizes the full procedure.

Algorithm 1 SCORE_OBN: score-difference function for a candidate move

Require: Current DAG $\mathcal{G}$; candidate move $m \in \{\text{add}, \text{delete}, \text{reverse}\}$; data $(\mathbf{X}, \mathbf{w}, \mathbf{Z})$; effective sample size n_{eff} ; scoring choice `score_type`

Ensure: Score difference Δ for applying m to $\mathcal{G}$

```

1: Let  $\mathcal{G}'$  denote the DAG obtained by applying  $m$  to  $\mathcal{G}$ 
2: if score_type = SPLIT-BIC and  $m$  = reverse then
3:    $\Delta \leftarrow \tilde{\ell}(\hat{\theta}_{\mathcal{G}'}^w | \mathcal{G}') - \tilde{\ell}(\hat{\theta}_{\mathcal{G}}^w | \mathcal{G})$  {Unpenalized rescaled pseudo-log-likelihood}
4: else
5:    $\Delta \leftarrow -\frac{1}{2} [\widetilde{\text{BIC}}(\mathcal{G}') - \widetilde{\text{BIC}}(\mathcal{G})]$  {Equation 7 of main text}
6: end if
7: return  $\Delta$ 

```

2. Proofs of Theoretical Results. Throughout, we use the notation established in Section 5 of the main text: ξ denotes the superpopulation, U_N the finite population of size N , S_n the probability sample of size n , $w_i = 1/\pi_i$ the survey weights, and $P_{\xi}(\mathbf{X} | \mathcal{G}, \theta, \mathbf{Z})$ the ordinal BN model under DAG $\mathcal{G}$ with parameters θ and covariates $\mathbf{Z}$. We define:

- Census log-likelihood: $\ell_U(\theta | \mathcal{G}) = \sum_{i \in U_N} \log P(\mathbf{x}_i | \mathcal{G}, \theta, \mathbf{z}_i)$.

Algorithm 2 swa-oBn: Survey-weighted, covariate-adjusted ordinal causal discovery

Require: Ordinal data $\mathbf{X} \in \mathbb{R}^{n \times q}$; survey weights $\mathbf{w} = (w_1, \dots, w_n)$ (optional); covariates $\mathbf{Z} \in \mathbb{R}^{n \times r}$ (optional); blacklist $\mathcal{B}$, whitelist $\mathcal{W}$ (optional); number of restarts R ; number of bootstrap replicates B ; scoring choice $\text{score_type} \in \{\text{SINGLE-BIC}, \text{SPLIT-BIC}\}$

Ensure: Point-estimate DAG $\hat{\mathcal{G}}$; edge inclusion probabilities $\{\hat{p}_{jk}\}$

{Preprocessing}
 1: Normalize weights: $w_i \leftarrow w_i \cdot n / \sum_{i'} w_{i'}$
 2: Compute $n_{\text{eff}} \leftarrow (\sum_i w_i)^2 / \sum_i w_i^2$
{Point estimate via multi-start hill-climbing}
 3: **for** $s = 1, \dots, R$ **do**
 4: Initialize $\mathcal{G}^{(s)}$ (empty graph if $s = 1$; random sparse DAG otherwise)
 5: Enforce whitelist edges in $\mathcal{G}^{(s)}$
 6: **repeat**
 7: For each candidate move m consistent with $\mathcal{B}$, $\mathcal{W}$, and acyclicity, compute $\Delta_m \leftarrow \text{SCORE_OBN}(\mathcal{G}^{(s)}, m, \mathbf{X}, \mathbf{w}, \mathbf{Z}, n_{\text{eff}}, \text{score_type})$
 8: Apply the move with the largest $\Delta_m > 0$, if any exists
 9: **until** no improving move exists
 10: **end for**
 11: $\hat{\mathcal{G}} \leftarrow \arg \min_s \widehat{\text{BIC}}(\mathcal{G}^{(s)})$
{Bootstrap uncertainty quantification}
 12: **for** $b = 1, \dots, B$ **do**
 13: Draw indices $\mathcal{I}^{(b)}$ by sampling $\{1, \dots, n\}$ with replacement
 14: Construct $(\mathbf{X}^{(b)}, \mathbf{w}^{(b)}, \mathbf{Z}^{(b)})$ from $\mathcal{I}^{(b)}$; recompute $n_{\text{eff}}^{(b)}$
 15: Run lines 3–10 on $(\mathbf{X}^{(b)}, \mathbf{w}^{(b)}, \mathbf{Z}^{(b)})$ to obtain $\hat{\mathcal{G}}^{(b)}$
 16: **end for**
 17: $\hat{p}_{jk} \leftarrow B^{-1} \sum_{b=1}^B \mathbf{1}(j \rightarrow k \in \hat{\mathcal{G}}^{(b)})$ for all (j, k)

- Pseudo-log-likelihood: $\tilde{\ell}(\theta | \mathcal{G}) = \sum_{i \in S_n} w_i \log P(\mathbf{x}_i | \mathcal{G}, \theta, \mathbf{z}_i)$.
- Superpopulation expected log-likelihood: $Q_\xi(\theta | \mathcal{G}) = \mathbb{E}_\xi[\log P(\mathbf{X} | \mathcal{G}, \theta, \mathbf{Z})]$.
- Normalized versions: $\bar{\ell}_U(\theta | \mathcal{G}) = N^{-1} \ell_U(\theta | \mathcal{G})$, $\tilde{\bar{\ell}}(\theta | \mathcal{G}) = N^{-1} \tilde{\ell}(\theta | \mathcal{G})$.

We use $\theta_{\mathcal{G}}^*$ to denote the superpopulation maximizer of $Q_\xi(\cdot | \mathcal{G})$, i.e., $\theta_{\mathcal{G}}^* = \arg \max_{\theta} Q_\xi(\theta | \mathcal{G})$; when $\mathcal{G} = \mathcal{G}^*$, this coincides with the true data-generating parameter, which we denote simply by θ^* . For a misspecified DAG $\mathcal{G}' \neq \mathcal{G}^*$, we use $\theta_{\mathcal{G}'}^\dagger \equiv \theta_{\mathcal{G}'}^*$ to emphasize its role as the KL-projection rather than a truth-level parameter. The *conditional* (covariate-pointwise) KL-projection is denoted $\theta^\dagger(\mathbf{z}) = \arg \max_{\theta} \sum_{\mathbf{x}} P(\mathbf{x} | \mathcal{G}^*, \theta^*, \mathbf{z}) \log P(\mathbf{x} | \mathcal{G}', \theta, \mathbf{z})$, which in general differs from the unconditional $\theta_{\mathcal{G}'}^\dagger$.

2.1. *Proof of Proposition 1 (Identifiability under Covariate Adjustment).* Let $\mathcal{G}^* \neq \mathcal{G}'$ be Markov-equivalent DAGs. Write $\vartheta = (\theta, \delta) \in \Theta \times \mathbb{R}^{rq}$ for the full parameter under the covariate-adjusted model, and define the pointwise pseudo-true parameter

$$\theta^\dagger(\mathbf{z}; \vartheta^*) = \arg \max_{\theta} \sum_{\mathbf{x}} P(\mathbf{x} | \mathcal{G}^*, \vartheta^*, \mathbf{z}) \log P(\mathbf{x} | \mathcal{G}', \theta, \mathbf{z}).$$

Define the exceptional set

$$\mathcal{E} = \left\{ \vartheta^* : \mathbb{E}_{\mathbf{Z}} \left[\text{KL} \left(P_\xi(\mathbf{X} | \mathcal{G}^*, \vartheta^*, \mathbf{Z}) \parallel P(\mathbf{X} | \mathcal{G}', \theta^\dagger(\mathbf{Z}; \vartheta^*), \mathbf{Z}) \right) \right] = 0 \right\}.$$

We show that $\mathcal{E}$ has Lebesgue measure zero in $\Theta \times \mathbb{R}^{rq}$.

2.1.1. *Step 1: Reduction to the no-covariate case.* For each fixed $\mathbf{Z} = \mathbf{z}$, the covariate-adjusted cumulative link model becomes

$$P_\xi(X_j \leq \ell \mid \mathbf{X}_{\text{pa}(j)}, \mathbf{Z} = \mathbf{z}) = F\left(\gamma_{j\ell} - \sum_{k \in \text{pa}(j)} \beta_{jk} X_k - \tilde{\alpha}_j(\mathbf{z})\right),$$

with $\tilde{\alpha}_j(\mathbf{z}) = \alpha_j + \boldsymbol{\delta}'_j \mathbf{z}$. This is a no-covariate oBN with thresholds $\{\gamma_{j\ell}\}$, parent-effect parameters $\{\beta_{jk}\}$, and shifted intercepts $\{\tilde{\alpha}_j(\mathbf{z})\}$.

2.1.2. *Step 2: The identifiability variety is intercept-invariant.* The identifiability proof of Ni and Mallick (2022) establishes that the ordinal parameterizations $P(X_j \mid X_k)$ and $P(X_k \mid X_j)$ are generically non-equivalent by analyzing a polynomial system in the threshold and parent-effect parameters. The intercept α_j determines the marginal location of the conditional distribution but does not enter the polynomial relations that govern distributional equivalence. Consequently, the exceptional set

$$\mathcal{V} = \left\{(\gamma, \beta, \alpha) : P(\mathbf{X} \mid \mathcal{G}^*, \gamma, \beta, \alpha) \equiv P(\mathbf{X} \mid \mathcal{G}', \theta_{\text{NC}}^\dagger(\gamma, \beta, \alpha))\right\}$$

in the no-covariate parameter space is a proper algebraic subvariety of Lebesgue measure zero, invariant under intercept translation. Here $\theta_{\text{NC}}^\dagger$ denotes the no-covariate pseudo-true parameter. By Assumption 5, $\mathcal{V}$ has Lebesgue measure zero.

2.1.3. *Step 3: The exceptional set has measure zero in $(\theta^*, \boldsymbol{\delta}^*)$ -space.* *Case (a): $\mathbf{Z}$ is continuous with a density on $\mathbb{R}^r$.* Fix $(\theta^*, \boldsymbol{\delta}^*)$ and consider the map $\mathbf{z} \mapsto (\gamma^*, \beta^*, \tilde{\alpha}^*(\mathbf{z}))$. This is an affine map. Since $\mathcal{V}$ has Lebesgue measure zero and the affine map has nonzero Jacobian for $\boldsymbol{\delta}^*$ outside a measure-zero degenerate set, the pre-image $\{\mathbf{z} : (\gamma^*, \beta^*, \tilde{\alpha}^*(\mathbf{z})) \in \mathcal{V}\}$ has Lebesgue measure zero in $\mathbf{z}$ -space. Therefore the pointwise KL is strictly positive for $P_{\mathbf{Z}}$ -almost every $\mathbf{z}$, and

$$\mathbb{E}_{\mathbf{Z}} \left[\text{KL} \left(P_\xi(\cdot \mid \mathcal{G}^*, \vartheta^*, \mathbf{Z}) \parallel P(\cdot \mid \mathcal{G}', \theta^\dagger(\mathbf{Z}; \vartheta^*), \mathbf{Z}) \right) \right] > 0$$

whenever $\boldsymbol{\delta}^*$ avoids this degenerate set. *Case (b): $\mathbf{Z}$ is discrete with finite support $\{\mathbf{z}_1, \dots, \mathbf{z}_K\}$.* For each k , define

$$\mathcal{E}_k = \left\{(\theta^*, \boldsymbol{\delta}^*) : (\gamma^*, \beta^*, \alpha_j^* + (\boldsymbol{\delta}_j^*)^\top \mathbf{z}_k) \in \mathcal{V}\right\}.$$

For each fixed $\mathbf{z}_k$, the map $(\theta^*, \boldsymbol{\delta}^*) \mapsto (\gamma^*, \beta^*, \alpha_j^* + (\boldsymbol{\delta}_j^*)^\top \mathbf{z}_k)$ is affine with nonzero Jacobian (provided $\mathbf{z}_k \neq \mathbf{0}$). By Fubini's theorem and the measure-zero property of $\mathcal{V}$, $\mathcal{E}_k$ has Lebesgue measure zero in $(\theta^*, \boldsymbol{\delta}^*)$ -space. The overall exceptional set satisfies $\mathcal{E} \subseteq \bigcup_{k=1}^K \mathcal{E}_k$, which, as a finite union of measure-zero sets, has Lebesgue measure zero. (When $\mathbf{z}_k = \mathbf{0}$ for some k , we redefine the origin so that all $\mathbf{z}_k$ are nonzero; this does not change the model.) Therefore $\mathcal{E}$ has Lebesgue measure zero, as claimed. When the true parameter $(\theta^*, \boldsymbol{\delta}^*) \notin \mathcal{E}$, the expected KL divergence is strictly positive. $\square$

2.2. *Proof of Proposition 2 (Pseudo-MLE Consistency).* We prove that $\hat{\theta}_{\mathcal{G}}^w \xrightarrow{p} \theta_{\mathcal{G}}^*$ under the joint $(\xi \times p_N)$ -probability. The argument proceeds in three steps.

2.2.1. *Step 1: Design consistency of the normalized pseudo-log-likelihood.* Fix $\theta \in \Theta_{\mathcal{G}}$ and define $g_{\theta}(\mathbf{x}, \mathbf{z}) = \log P(\mathbf{x} \mid \mathcal{G}, \theta, \mathbf{z})$. Since $\mathbf{X}$ takes values in the finite set $\prod_{j=1}^q \{1, \dots, L_j\}$, Assumption 6 (Non-Degeneracy) gives $P(\mathbf{x} \mid \mathcal{G}, \theta, \mathbf{z}) > 0$ for all $(\mathbf{x}, \mathbf{z})$, and compactness of $\Theta_{\mathcal{G}}$ (Assumption 10(i)) combined with continuity of F yields a uniform lower bound $P(\mathbf{x} \mid \mathcal{G}, \theta, \mathbf{z}) \geq p_{\min} > 0$ over all $(\mathbf{x}, \mathbf{z}, \theta)$. Therefore g_{θ} is bounded uniformly in $\theta \in \Theta_{\mathcal{G}}$. By Assumption 9 (Design Consistency), applied to the function g_{θ} of $(\mathbf{x}, \mathbf{z})$,

$$(1) \quad \bar{\ell}(\theta \mid \mathcal{G}) - \bar{\ell}_U(\theta \mid \mathcal{G}) = \frac{1}{N} \sum_{i \in S_n} w_i g_{\theta}(\mathbf{x}_i, \mathbf{z}_i) - \frac{1}{N} \sum_{i \in U_N} g_{\theta}(\mathbf{x}_i, \mathbf{z}_i) \xrightarrow{p_N} 0.$$

Design Consistency in turn follows from Horvitz–Thompson consistency under Assumptions 7–8 (Positivity, Bounded Weights); see [Fuller \(2009\)](#), Ch. 1.

2.2.2. *Step 2: Superpopulation law of large numbers for the census log-likelihood.* By the i.i.d. superpopulation structure, $\{(\mathbf{x}_i, \mathbf{z}_i)\}_{i \in U_N}$ are i.i.d. draws from ξ . Since g_{θ} is bounded, the strong law of large numbers gives

$$(2) \quad \bar{\ell}_U(\theta \mid \mathcal{G}) = \frac{1}{N} \sum_{i=1}^N g_{\theta}(\mathbf{x}_i, \mathbf{z}_i) \xrightarrow{\text{a.s.}} \mathbb{E}_{\xi}[g_{\theta}(\mathbf{X}, \mathbf{Z})] = Q_{\xi}(\theta \mid \mathcal{G}).$$

2.2.3. *Step 3: Uniform convergence and argmax convergence.* Combining (1) and (2), for each fixed θ : $\bar{\ell}(\theta \mid \mathcal{G}) \xrightarrow{p} Q_{\xi}(\theta \mid \mathcal{G})$. We upgrade this to uniform convergence over $\Theta_{\mathcal{G}}$. By Assumption 10(iv), for any $\theta, \theta' \in \Theta_{\mathcal{G}}$,

$$|g_{\theta}(\mathbf{x}, \mathbf{z}) - g_{\theta'}(\mathbf{x}, \mathbf{z})| \leq M \|\theta - \theta'\|$$

for some constant $M > 0$. This Lipschitz condition, combined with the compactness of $\Theta_{\mathcal{G}}$ (Assumption 10(i)), gives equicontinuity. Pointwise convergence plus equicontinuity on a compact set gives uniform convergence ([van der Vaart, 1998](#), Theorem 2.4):

$$(3) \quad \sup_{\theta \in \Theta_{\mathcal{G}}} \left| \bar{\ell}(\theta \mid \mathcal{G}) - Q_{\xi}(\theta \mid \mathcal{G}) \right| \xrightarrow{p} 0.$$

By Assumption 10(iii), $Q_{\xi}(\theta \mid \mathcal{G})$ has a unique maximizer $\theta_{\mathcal{G}}^*$ with nonsingular Hessian. By (3) and the argmax continuous mapping theorem ([van der Vaart, 1998](#), Theorem 5.7),

$$(4) \quad \hat{\theta}_{\mathcal{G}}^w \xrightarrow{p} \theta_{\mathcal{G}}^*.$$

As a corollary:

$$(5) \quad \bar{\ell}(\hat{\theta}_{\mathcal{G}}^w \mid \mathcal{G}) \xrightarrow{p} Q_{\xi}(\theta_{\mathcal{G}}^* \mid \mathcal{G}).$$

□

2.3. *Proof of Theorem 5.1 (Skeleton Consistency).* Let $\mathcal{G}^*$ be the true DAG and $\mathcal{G}'$ a DAG not Markov-equivalent to $\mathcal{G}^*$. We show $\Pr(\widehat{\text{BIC}}(\mathcal{G}^*) < \widehat{\text{BIC}}(\mathcal{G}')) \rightarrow 1$.

Write

$$(6) \quad \widehat{\text{BIC}}(\mathcal{G}^*) - \widehat{\text{BIC}}(\mathcal{G}') = -2 \left[\bar{\ell}(\hat{\theta}_{\mathcal{G}^*}^w \mid \mathcal{G}^*) - \bar{\ell}(\hat{\theta}_{\mathcal{G}'}^w \mid \mathcal{G}') \right] + (K_{\mathcal{G}^*} - K_{\mathcal{G}'})\rho_n.$$

2.3.1. *Case 1A: $\mathcal{G}'$ is an overfit model ($\mathcal{G}^* \subset \mathcal{G}'$ as directed graphs, so $K_{\mathcal{G}'} > K_{\mathcal{G}^*}$).* Since $\mathcal{G}^*$ is a subgraph of $\mathcal{G}'$, the oBN model under $\mathcal{G}^*$ is nested within the model under $\mathcal{G}'$, and both achieve the same superpopulation optimum: $Q_\xi(\theta_{\mathcal{G}^*}^* | \mathcal{G}^*) = Q_\xi(\theta_{\mathcal{G}'}^\dagger | \mathcal{G}')$. By the nested-model likelihood-ratio theory for pseudo-maximum likelihood (Rao and Scott, 1981; Lumley and Scott, 2014), the survey-weighted pseudo-log-likelihood ratio is

$$2 \left[\tilde{\ell}(\hat{\theta}_{\mathcal{G}'}^w | \mathcal{G}') - \tilde{\ell}(\hat{\theta}_{\mathcal{G}^*}^w | \mathcal{G}^*) \right] = O_p(1),$$

asymptotically distributed as a weighted mixture of chi-squared random variables with weights determined by the design effects of the added parameters. Meanwhile, $(K_{\mathcal{G}^*} - K_{\mathcal{G}'})\rho_n < 0$ with $|K_{\mathcal{G}^*} - K_{\mathcal{G}'}|\rho_n \rightarrow \infty$ since $\rho_n \rightarrow \infty$. Therefore

$$\widetilde{\text{BIC}}(\mathcal{G}^*) - \widetilde{\text{BIC}}(\mathcal{G}') = O_p(1) + (K_{\mathcal{G}^*} - K_{\mathcal{G}'})\rho_n \rightarrow -\infty$$

in probability.

2.3.2. *Case 1B: $\mathcal{G}'$ has more edges than $\mathcal{G}^*$ but is not a supergraph.* In this case, $\mathcal{G}'$ is misspecified: it both adds spurious edges and omits at least one true edge. The omission places $\mathcal{G}'$ in the regime of Case 2 below, where the likelihood gap dominates the penalty difference regardless of sign.

2.3.3. *Case 2: $\mathcal{G}'$ is an underfit model (missing edges from $\mathcal{G}^*$).* By Assumption 3 (Faithfulness), every edge in $\mathcal{G}^*$ corresponds to a genuine conditional dependence: $Q_\xi(\theta^* | \mathcal{G}^*) > Q_\xi(\theta^\dagger | \mathcal{G}')$. Let $\Delta = Q_\xi(\theta^* | \mathcal{G}^*) - Q_\xi(\theta^\dagger | \mathcal{G}') > 0$. By (5), the likelihood term is $N(\Delta + o_p(1))$, dominating the $O(\rho_n)$ penalty since $\rho_n/N \rightarrow 0$:

$$\widetilde{\text{BIC}}(\mathcal{G}^*) - \widetilde{\text{BIC}}(\mathcal{G}') = -2N(\Delta + o_p(1)) + O(\rho_n) \rightarrow -\infty.$$

2.3.4. *Case 3: Same skeleton, not Markov-equivalent.* This can only occur if $\mathcal{G}'$ and $\mathcal{G}^*$ differ in v-structures. By Assumption 3, $\mathcal{G}'$ implies at least one conditional independence not present in P_ξ , making it effectively misspecified (as in Case 2). $\square$

2.4. *Proof of Theorem 5.2 (Orientation Consistency).* Let $\mathcal{G}^*$ be the true DAG and $\mathcal{G}'$ a Markov-equivalent DAG with $\mathcal{G}' \neq \mathcal{G}^*$. We show $\Pr(\tilde{\ell}(\hat{\theta}_{\mathcal{G}^*}^w | \mathcal{G}^*) > \tilde{\ell}(\hat{\theta}_{\mathcal{G}'}^w | \mathcal{G}')) \rightarrow 1$.

2.4.1. *Step 1: Convergence of the pseudo-log-likelihood ratio.* Applying (5) to both $\mathcal{G}^*$ and $\mathcal{G}'$:

$$(7) \quad \frac{1}{N} \left[\tilde{\ell}(\hat{\theta}_{\mathcal{G}^*}^w | \mathcal{G}^*) - \tilde{\ell}(\hat{\theta}_{\mathcal{G}'}^w | \mathcal{G}') \right] \xrightarrow{p} Q_\xi(\theta_{\mathcal{G}^*}^* | \mathcal{G}^*) - Q_\xi(\theta_{\mathcal{G}'}^\dagger | \mathcal{G}'),$$

where $\theta_{\mathcal{G}^*}^* = \theta^*$ (since $\mathcal{G}^*$ is correctly specified) and $\theta_{\mathcal{G}'}^\dagger = \arg \max_\theta Q_\xi(\theta | \mathcal{G}')$ is the superpopulation KL-projection onto the model under $\mathcal{G}'$.

2.4.2. *Step 2: Identification of the limit as an expected KL divergence.* Expanding the definition of Q_ξ :

$$\begin{aligned} Q_\xi(\theta^* | \mathcal{G}^*) - Q_\xi(\theta_{\mathcal{G}'}^\dagger | \mathcal{G}') &= \mathbb{E}_\xi \left[\log P(\mathbf{X} | \mathcal{G}^*, \theta^*, \mathbf{Z}) - \log P(\mathbf{X} | \mathcal{G}', \theta_{\mathcal{G}'}^\dagger, \mathbf{Z}) \right] \\ &= \mathbb{E}_\mathbf{Z} \left\{ \mathbb{E}_\xi \left[\log \frac{P(\mathbf{X} | \mathcal{G}^*, \theta^*, \mathbf{Z})}{P(\mathbf{X} | \mathcal{G}', \theta_{\mathcal{G}'}^\dagger, \mathbf{Z})} \middle| \mathbf{Z} \right] \right\} \\ &\geq \mathbb{E}_\mathbf{Z} \left\{ \mathbb{E}_\xi \left[\log \frac{P(\mathbf{X} | \mathcal{G}^*, \theta^*, \mathbf{Z})}{P(\mathbf{X} | \mathcal{G}', \theta^\dagger(\mathbf{Z}), \mathbf{Z})} \middle| \mathbf{Z} \right] \right\} \\ &= \mathbb{E}_\mathbf{Z} \left[\text{KL} \left(P_\xi(\mathbf{X} | \mathcal{G}^*, \theta^*, \mathbf{Z}) \parallel P(\mathbf{X} | \mathcal{G}', \theta^\dagger(\mathbf{Z}), \mathbf{Z}) \right) \right], \end{aligned} \quad (8)$$

where the inequality in (8) holds because $\theta^\dagger(z)$ is the pointwise KL-minimizer given $\mathbf{Z} = z$, so each inner expectation is at least as large as when evaluated at any other parameter. By Proposition 1, the final expected KL is strictly positive under the generic-parameter assumption on (θ^*, δ^*) : let $\Delta > 0$ denote its value.

2.4.3. Step 3: Conclusion. The pseudo-log-likelihood difference is $N[\Delta + o_p(1)] \rightarrow +\infty$ in probability, establishing the result. $\square$

REMARK 1. *No penalty term appears in this proof. The orientation comparison is a pseudo-likelihood ratio, and its consistency follows from generic ordinal identifiability (Proposition 1) and pseudo-MLE consistency (Proposition 2). Both scoring variants of Section 4 inherit this consistency: split-BIC compares reversals on the pseudo-likelihood directly, while single-BIC adds a penalty differential of order $O(\rho_n)$ that is asymptotically dominated by the $O(N)$ likelihood gap. The variants differ only in finite-sample behavior under heterogeneous $\mathbf{L}$; see Section 6.*

3. Polychoric Clustering Details. Figure 1 presents the weighted polychoric correlation matrix and hierarchical clustering dendrogram for the 13 SHEP social need items. Items within each highlighted cluster have pairwise polychoric correlations $\rho \geq 0.75$.

4. Simulation Study 1: Edge Precision and Orientation Accuracy. Figure 2 presents edge precision (TP/(TP+FP)) for the `swa-oBn` versus naive `oBn` across design strength and signal strength conditions from Study 1 (Section 6.2 of the main text). Under mild and strong designs, the `swa-oBn` achieves substantially higher precision than the naive `oBn` (approximately 0.90–0.95 versus 0.70–0.75), indicating that the primary benefit of survey weighting is the elimination of spurious edges. Under no design, precision is identical for both methods.

Figure 3 presents the correct orientation rate (fraction of true edges found and correctly oriented). The weighted estimator records modestly lower orientation rates than the naive estimator under informative sampling, reflecting a precision-recall tradeoff: by suppressing spurious edges (Figure 2), `swa-oBn` also retains fewer accidentally correctly oriented edges. The combined effect on SHD is favorable, as documented in Section 6.2 of the main text.

5. Simulation Study 3 Supplementary Tables. This appendix reports the full 6×4 grid of paired SHD gaps (split-BIC – single-BIC) underlying the headline result of Section 6.4. Tables 1 and 2 cover the weighted and unweighted arms respectively. Negative values indicate that split-BIC achieves lower SHD; significance is assessed at $|\text{gap}| > 1.96 \cdot \text{SE}$, with significant cells marked *.

The taxonomy of Table 1 is qualitatively stable across n : split-BIC dominates in five regimes and single-BIC dominates under *Graded*, with the same sign pattern at every sample size. Gap magnitudes grow with n , consistent with the per-edge bias of the unpenalized second pass accumulating with effective sample size. The unweighted arm (Table 2) exhibits the same regime structure with slightly larger gap magnitudes. The choice between single-BIC and split-BIC scoring thus depends on the $\mathbf{L}$ -distribution alone, independent of survey weighting; the regime taxonomy of Section 6.4 applies to both `swa-oBn` and the unweighted `oBn` of Ni, Chen and Wang (2025).

6. Simulation Study 4: Survey Weighting for Other Algorithm Families. To demonstrate that the distortion introduced by informative sampling is not specific to ordinal Bayesian networks, we evaluate survey-weighted and unweighted versions of three additional causal discovery algorithms: the PC algorithm (Spirtes, Glymour and Scheines, 2000), the FCI algorithm (Spirtes, Glymour and Scheines, 2000), and LiNGAM (Shimizu et al., 2006).




FIG 1. Polychoric correlation matrix and hierarchical clustering of 13 SHEP social need items. Shaded blocks indicate clusters with $\rho \geq 0.75$: basic needs (food, housing, paying for basics), work/education, and isolation/loneliness/discrimination.

TABLE 1

Study 3, weighted arm: paired SHD gap (split-BIC – single-BIC) with standard errors across six $\mathbf{L}$ -regimes and four sample sizes. Negative values favor split-BIC; * denotes $|\text{gap}| > 1.96 \cdot \text{SE}$.

Regime	$\mathbf{L}$ -set	$n = 200$	$n = 400$	$n = 700$	$n = 1,000$
Uniform	{3}	−0.07 (0.03)*	−0.05 (0.03)	−0.13 (0.06)*	−0.07 (0.02)*
Adjacent	{3, 4}	−0.23 (0.05)*	−0.23 (0.06)*	−0.28 (0.06)*	−0.46 (0.06)*
Graded	{3, 4, 5}	+0.07 (0.05)	+0.29 (0.05)*	+0.47 (0.04)*	+0.42 (0.04)*
Polarized	{3, 5}	−0.21 (0.04)*	−0.24 (0.05)*	−0.36 (0.05)*	−0.61 (0.05)*
Wide	{3, 7}	−0.19 (0.04)*	−0.19 (0.04)*	−0.61 (0.05)*	−0.98 (0.07)*
Extreme	{3, 9}	−0.20 (0.04)*	−0.31 (0.05)*	−0.52 (0.05)*	−0.81 (0.06)*

6.1. *Setup.* We generate a population of $N = 20,000$ units from a three-variable linear chain $X \rightarrow Y \rightarrow Z$ with continuous variables ($X, Z \sim N(0, 1)$, $Y = 0.8X + \varepsilon_Y$, $Z = 0.8Y + \varepsilon_Z$, $\varepsilon_Y, \varepsilon_Z \sim N(0, 1)$). From this population, Bernoulli samples are drawn under three design conditions (none, mild, strong), with oversampling based on the top quartile of the combined score of X and Z . Each scenario is replicated 100 times.




FIG 2. Edge precision ($TP/(TP+FP)$) of *swa-oBN* versus naive oBN across design strength (left) and signal strength (right) conditions. $q = 8$ nodes, 100 replications.




FIG 3. Correct orientation rate (fraction of true edges found and correctly oriented) across design strength (left) and signal strength (right) conditions. $q = 8$ nodes, 100 replications.

For the PC and FCI algorithms, survey weighting is implemented by replacing standard conditional independence tests with design-adjusted Wald tests using the `survey` package in R (Lumley, 2004). For LiNGAM, we implement a weighted pseudo-population resampling procedure that restores i.i.d. structure by resampling from the observed data proportional to survey weights prior to applying the standard ICA-based algorithm.

TABLE 2

Study 3, unweighted arm: paired SHD gap (split-BIC – single-BIC) with standard errors across six $\mathcal{L}$ -regimes and four sample sizes. Negative values favor split-BIC; * denotes $|\text{gap}| > 1.96 \cdot \text{SE}$.

Regime	$\mathcal{L}$ -set	$n = 200$	$n = 400$	$n = 700$	$n = 1,000$
Uniform	{3}	−0.08 (0.04)	−0.15 (0.04)*	−0.05 (0.04)	−0.16 (0.03)*
Adjacent	{3, 4}	−0.27 (0.05)*	−0.38 (0.06)*	−0.50 (0.06)*	−0.75 (0.05)*
Graded	{3, 4, 5}	+0.10 (0.05)*	+0.06 (0.04)	+0.36 (0.04)*	+0.15 (0.04)*
Polarized	{3, 5}	−0.19 (0.06)*	−0.52 (0.05)*	−0.71 (0.05)*	−0.74 (0.05)*
Wide	{3, 7}	−0.31 (0.04)*	−0.49 (0.05)*	−1.03 (0.05)*	−1.06 (0.06)*
Extreme	{3, 9}	−0.19 (0.05)*	−0.55 (0.04)*	−0.65 (0.06)*	−1.03 (0.06)*

6.2. *Results.* Figure 4 presents the correct skeleton recovery rate across the three design conditions for all three algorithm families. Under simple random sampling (design = none), all six estimators (three algorithms $\times$ two weighting conditions) recover the correct skeleton in nearly all replications. Under mild and strong designs, the unweighted versions show a marked decline in skeleton recovery, while the survey-weighted versions maintain substantially higher correct recovery rates.

Figure 5 presents the spurious edge rate (proportion of replications in which a direct $X-Z$ edge is inferred). Under no design, spurious edges are rare for all methods. Under mild and strong designs, the unweighted versions infer the spurious $X-Z$ edge at substantially elevated rates relative to their weighted counterparts. The effect is largest for FCI (95% spurious-edge rate unweighted vs. 4% weighted under strong design), substantial for the PC algorithm, and present but smaller for LiNGAM, whose ICA-based identification is less sensitive to design-induced marginal dependence.

These results confirm that the survey-weighting problem documented in Studies 1 and 2 for the ordinal Bayesian network is a general phenomenon affecting all three major families of causal discovery algorithms.




FIG 4. Study 4: Correct skeleton recovery rate for survey-weighted and unweighted versions of PC, FCI, and LiNGAM across three design conditions. True graph: $X \rightarrow Y \rightarrow Z$. 100 replications per scenario.




FIG 5. Study 4: Spurious X - Z edge rate for survey-weighted and unweighted versions of PC, FCI, and LiNGAM across three design conditions. True graph: $X \rightarrow Y \rightarrow Z$. 100 replications per scenario.

7. Proportional Odds Diagnostics. The `swa-obn` framework rests on the cumulative link specification of Assumption 4 of the main text, of which the proportional odds (PO) constraint is the most consequential working restriction: a single set of slope parameters is shared across the $L_j - 1$ cumulative thresholds for node j . This appendix reports diagnostic checks of PO for the application of Section 7 of the main text, documents the limitations of off-the-shelf diagnostics under a complex survey design, and presents a design-based sensitivity check for the regression whose orientation most directly drives our substantive conclusions.

7.1. Limitations of off-the-shelf diagnostics under a complex survey. The standard PO diagnostic is the [Brant \(1990\)](#) score test, implemented in R via `brant : : brant`. The test fits the cumulative link model via maximum likelihood and rejects PO if the between-threshold slope differences are jointly too large relative to the estimated covariance. In survey settings this procedure has two well-known failings. First, the estimator is a pseudo-MLE that treats observations as independently and identically distributed; the Hessian used to construct the score-test variance therefore omits the between-stratum and unequal-probability components of design variance and is anticonservative. Under design effect $deff$, the nominal test statistic is inflated by a factor of approximately $deff$, and at our sample size ($n_{eff} = 1,387$, $deff = 3.89$) this inflation can transform a modest PO departure into an apparently decisive rejection. Second, all B bootstrap replicates of the `swa-obn` refit cumulative link regressions at parent sets that are themselves estimated; the nominal Brant distribution does not account for this parent-set stochasticity. We therefore present nominal Brant results as indicative rather than confirmatory, and follow them with a design-based analog for the regression in the discovered DAG whose orientation drives our principal finding.

7.2. Nominal Brant results. We report Brant tests under two parent specifications. *Pre-analysis* regresses each of the $q = 10$ ordinal nodes on the demographic covariates $\mathbf{Z}$ alone (age band and race-ethnicity-sex), a specification that makes no use of the discovered DAG. *Post-analysis* regresses each node on its estimated parent set from the consensus DAG (Figure 3 of the main text) together with $\mathbf{Z}$. At nominal $\alpha = 0.05$, the pre-analysis Brant test rejects PO in 1 of the 10 nodes; the post-analysis test rejects PO in 3 of 10. The increase in rejections from pre- to post-analysis is consistent with the Brant statistic inflating mechanically as the parent set grows (each added parent contributes $L_k - 1$ additional slopes to the joint test) and does not, on its own, indicate that the PO structure deteriorates conditional on the discovered parents. Given the anticonservatism documented in Section 7.1, we do not interpret these counts as evidence against the working model.

7.3. Working-model interpretation. Following Lumley and Scott (2015), we treat the PO cumulative link as an explicitly designated working model under complex-survey pseudo-likelihood. The generic identifiability result invoked throughout this paper (Proposition 1 of the main text) is a statement about the ordinal parameterization itself, not about whether the data-generating mechanism exactly satisfies PO. Under the working-model view, the pseudo-MLE $\hat{\theta}^w$ consistently estimates the KL projection of the population distribution onto the PO model under the sampling design; the bootstrap inclusion probabilities $\hat{p}_{jk}$ describe the dependency structure of this projection, not of the unrestricted population. The substantive consequence is that the consensus DAG should be read as summarizing the dominant directed dependencies in the SHEP data under the ordinal working model, which is the claim that the theory of Section 5 of the main text supports.

8. Survey Weight Distribution. Table 3 summarizes the distribution of survey weights in the SHEP analytic sample. The design effect ($\text{DEFF} = n/n_{\text{eff}}$) quantifies the loss of effective sample size due to unequal weighting.

TABLE 3
Survey weight distribution and design effect.

	n	n_{eff}	DEFF	Min	Median	Max	CV
Overall	5,396	1,387	3.89	3.4	58.4	3,330.1	1.700

9. Composite Node Construction. Table 4 maps the ten analytic nodes used in Section 7 of the main text to their constituent SHEP items, following the polychoric clustering of Appendix 3. Composites are constructed row-wise as the maximum reported severity across constituent items, preserving the three-level ordinal scale (no support needed; needed support and got it; needed support but did not get it). The resulting level vector $L = (3, 3, 3, 3, 3, 3, 3, 3, 4, 5)$ places the application in the *Graded* regime of Section 6.4, motivating the use of single-BIC scoring.

TABLE 4
Mapping of analytic nodes to constituent SHEP items. The first eight rows are need nodes; the last two are health outcomes.

Analytic node	Constituent SHEP items (Q43 subitem)	L_j	Construction
Caregiving	Obtaining adult caregiving for yourself or others	3	single
Childcare	Obtaining childcare	3	single
Transportation	Getting transportation for basic needs	3	single
Internet	Accessing the internet at home	3	single
Legal	Getting assistance with legal issues	3	single
Basic Needs	Paying for basics like food, housing, medical care and heating; paying for food; getting or maintaining housing	3	composite
Work/Education	Finding or keeping work; getting additional education or job training	3	composite
D/I	Feeling socially isolated; feeling lonely; managing experiences of discrimination	3	composite
Mental Health	Q45: overall mental or emotional health rating	4	single
Well-being	Q46–48: satisfaction, involvement, functioning	5	composite

10. Complete Bootstrap Edge Probability Matrix. Table 5 reports the full 10×10 matrix of bootstrap edge inclusion probabilities for both the weighted and unweighted analyses. Entry (j, k) gives $\hat{p}_{jk} = B^{-1} \sum_{b=1}^B \mathbf{1}(j \rightarrow k \in \hat{\mathcal{G}}^{(b)})$; rows index the parent (from) node and columns the child (to) node.

TABLE 5

Full 10×10 bootstrap edge inclusion probability matrix ($B = 250$ resamples). Each cell reports the weighted probability followed by the unweighted in brackets: $\hat{p}_{jk}^{wtd} [\hat{p}_{jk}^{unwtd}]$. Rows index the parent (from) node and columns the child (to) node, so cell (j, k) gives $\hat{p}_{j \rightarrow k}$. Diagonal cells (“—”) and the sixteen outcome-to-need cells forbidden by the blacklist (all $\hat{p} = 0$ in both specifications) are omitted. Boldface marks the fifteen directed cells whose weighted and unweighted inclusion probabilities fall on opposite sides of the 0.50 consensus threshold. The self-reported nodes in the table include: (1) caregiving (adult caregiving for self or others); (2) childcare; (3) transportation (for basic needs like medical care or grocery shopping); (4) internet (getting internet at home); (5) legal (obtaining assistance for legal issues); (6) basic needs (i.e., paying for basics like food, housing, medical care, and heating; paying for food (only); getting or maintaining housing); (7) work/education (finding or keeping work, getting additional education or job training); (8) isolation/loneliness/discrimination, abbreviated D/I (feeling socially isolated; feeling lonely; managing experiences of discrimination); (9) mental health; and (10) well-being.

[illegible]

11. Weighted versus Unweighted Comparison. Three differences sever the mediation pathway underlying the paper’s principal finding. First, the basic-needs–D/I edge reverses: weighted selects basic needs $\rightarrow$ D/I ($\hat{p} = 0.68$); unweighted selects the reverse, D/I $\rightarrow$ basic needs ($\hat{p} = 0.67$). Second, a direct basic-needs $\rightarrow$ well-being edge at $\hat{p} = 0.92$ unweighted falls to 0.15 weighted, spuriously bypassing the mediator. Third, the




FIG 6. Consensus DAGs with survey weighting (left) and without (right), on a shared node layout. Covariate adjustment, blacklist, restarts ($R = 10$), and bootstrap replicates ($B = 250$) are identical across panels; only weights differ. Edge styling matches main-text Figure 3 (solid: robust, $\hat{p} \geq 0.75$; dashed: suggestive, $0.50 \leq \hat{p} < 0.75$; thin grey: weak, $\hat{p} < 0.50$). The self-reported social need nodes in the figure include: (1) basic needs (i.e., paying for basics like food, housing, medical care, and heating; paying for food (only); getting or maintaining housing); (2) work/education (finding or keeping work, getting additional education or job training); (3) isolation/loneliness/discrimination, abbreviated D/I (feeling socially isolated; feeling lonely; managing experiences of discrimination); (4) caregiving (adult caregiving for self or others); (5) transportation (for basic needs like medical care or grocery shopping); (6) internet (getting internet at home); (7) legal (obtaining assistance for legal issues). Childcare was not connected to other nodes; thus, it is not shown in the figure.

internet–D/I edge reverses: weighted selects internet $\rightarrow$ D/I ($\hat{p} = 0.65$); unweighted selects D/I $\rightarrow$ internet ($\hat{p} = 0.95$). Together these produce the qualitatively different policy conclusion of Sections 7 and 8.

12. Causal Effect Estimation via the Recovered DAG. The bootstrap consensus DAG of Section 7 of the main text encodes a directed dependency structure among unmet social needs and Veteran mental health. Treated as a working causal model, this structure permits projection of intervention effects through the standard g-formula (Robins, 1986). We compute such projections for the single intervention most relevant to the paper’s policy narrative, $\text{do}(\text{BN} = \text{treated})$ — that is, resolving unmet basic needs — and report two complementary summaries: the total effect on mental health, and the controlled direct effect at a reference level of D/I. The contrast between these two summaries is the quantitative form of the bidirectional argument made in Section 8 of the main text and is the only result this section is intended to support.

12.1. Estimand and identification. Let Y denote mental health (five ordinal levels, 1 = best, 5 = worst), $A = \text{BN}$ denote unmet basic needs (three ordinal levels), $M = \text{D/I}$ denote isolation/loneliness/discrimination (three ordinal levels), and L denote the remaining nodes of the weighted DAG of Section 7. Under the assumed structure, the total effect of an intervention on A is identified by the g-formula

$$(9) \quad \mathbb{E}[Y \mid \text{do}(A = a)] = \sum_{m, \ell} \mathbb{E}[Y \mid A = a, M = m, L = \ell] \Pr(M = m, L = \ell \mid A = a),$$

and the controlled direct effect of A at $M = m^*$ by

$$(10) \quad \mathbb{E}[Y \mid \text{do}(A = a), \text{do}(M = m^*)] = \sum_{\ell} \mathbb{E}[Y \mid A = a, M = m^*, L = \ell] \Pr(L = \ell \mid A = a).$$

Identification rests on causal sufficiency conditional on Z and on the proportional-odds working assumption discussed; both are working approximations, and the resulting effects should be read as illustrative quantities under the assumed structure rather than as formal treatment-effect estimates.

12.2. Estimation procedure. Conditional expectations were estimated by survey-weighted ordinal logistic regression at each node, using `svyolr` from the `survey` package with a `polr` fallback for nodes at which `svyolr` failed to converge. For each of the $B = 250$ bootstrap resamples used to construct the consensus DAG, we re-fit the local conditional models on the same resample and evaluated (9)–(10) at $a = \text{‘needed support and got it’}$ relative to $a = \text{‘no support needed’}$, holding $m^* = \text{‘no support needed’}$ in the controlled direct effect. Bootstrap percentile intervals were computed across the 250 resamples. The bootstrap replay was seeded to reuse the exact resamples generated for structure learning, so structure and effect estimates share the same uncertainty budget.

12.3. Results and interpretation. The four projected quantities are reported in Table 6. Resolving unmet basic needs corresponds to an average mental-health improvement of 0.31 ordinal levels and a 4.9-percentage-point reduction in the probability of the worst mental-health response under the total-effect specification. The controlled direct effect at $\text{D/I} = \text{‘no support needed’}$ is markedly smaller: 1.0 percentage points. Approximately seventy percent of the projected mental-health benefit of resolving unmet basic needs is therefore eliminable by intervention on D/I, consistent with the structural finding that D/I sits as the dominant proximal mediator between basic needs and mental health.

The substantive reading of these numbers is that the recovered DAG is not merely consistent with a mediation interpretation — it is informative about the magnitude of mediation.

TABLE 6
Projected effects of $\text{do}(BN = \text{treated})$ under the recovered weighted DAG. Bootstrap percentile intervals ($B = 250$). $Y = \text{mental health}$.

Quantity	Estimate	95% bootstrap CI
Total effect on mean Y (ordinal levels)	+0.31	(0.13, 0.73)
Total effect on $\Pr(Y = \text{worst})$ (pp)	+4.9	(2.1, 10.8)
Controlled direct effect on $\Pr(Y = \text{worst})$ at D/I-ref (pp)	+1.0	(0.0, 10.0)
Eliminable share via D/I	≈ 0.70	—

Roughly seventy percent of the projected mental-health benefit of resolving basic needs operates through reductions in isolation/loneliness/discrimination; the remaining roughly thirty percent operates through other pathways in the network, including the caregiving→well-being→mental-health route identified in Section 7. This decomposition is the quantitative analog of the bidirectional intervention argument of Section 8: addressing basic needs without simultaneously addressing the mediator captures only the smaller direct share of the projected benefit, and addressing the mediator without addressing upstream basic needs leaves the larger upstream share unrealized.

We caution against over-interpreting these specific magnitudes. The projected effects depend on the recovered DAG, on the proportional-odds working assumption, and on causal sufficiency conditional on $\mathbf{Z}$; a misspecification of any of these would propagate to the numerical estimates while leaving the qualitative finding (that a substantial majority of the basic-needs effect operates through D/I) more robust. The wide bootstrap intervals — particularly for the controlled direct effect — reflect this combined uncertainty, and we read the seventy-percent figure as a coarse summary of mediation structure rather than as a precise effect-size estimate.

REFERENCES

- BRANT, R. (1990). Assessing Proportionality in the Proportional Odds Model for Ordinal Logistic Regression. *Biometrics* **46** 1171–1178.
- FULLER, W. A. (2009). *Sampling Statistics*. John Wiley & Sons, Hoboken, NJ.
- LUMLEY, T. (2004). Analysis of complex survey samples. *Journal of Statistical Software* **9** 1–19.
- LUMLEY, T. and SCOTT, A. (2014). Tests for regression models fitted to survey data. *Australian & New Zealand Journal of Statistics* **56** 1–14.
- LUMLEY, T. and SCOTT, A. (2015). AIC and BIC for Modeling with Complex Survey Data. *Journal of Survey Statistics and Methodology* **3** 1–18.
- NI, Y., CHEN, S. and WANG, Z. (2025). Causal Structural Modeling of Survey Questionnaires via a Bootstrapped Ordinal Bayesian Network Approach. *Psychometrika* **90** 229–250.
- NI, Y. and MALLICK, B. (2022). Ordinal Causal Discovery. In *Proceedings of the Thirty-Eighth Conference on Uncertainty in Artificial Intelligence* (J. CUSSENS and K. ZHANG, eds.). *Proceedings of Machine Learning Research* **180** 1530–1540. PMLR.
- RAO, J. N. K. and SCOTT, A. J. (1981). The analysis of categorical data from complex sample surveys: chi-squared tests for goodness of fit and independence in two-way tables. *Journal of the American Statistical Association* **76** 221–230.
- ROBINS, J. (1986). A new approach to causal inference in mortality studies with a sustained exposure period—application to control of the healthy worker survivor effect. *Mathematical modelling* **7** 1393–1512.
- SHIMIZU, S., HOYER, P. O., HYVÄRINEN, A. and KERMINEN, A. (2006). A Linear Non-Gaussian Acyclic Model for Causal Discovery. *Journal of Machine Learning Research* **7** 2003–2030.
- SPIRITES, P., GLYMOUR, C. N. and SCHEINES, R. (2000). *Causation, Prediction, and Search*, 2nd ed. MIT Press, Boston, MA.
- VAN DER VAART, A. W. (1998). *Asymptotic Statistics*. Cambridge University Press.